

TP1 - Optimizing Memory Access

Imad Kissami

October, 2025

Exercise 1

- This exercise aims to explore the impact of memory access strides on the performance of a C program.
- The following program allocates an array of doubles, initializes it to 1.0, and then performs a summation while traversing the array with different strides.

```
1  #include "stdio.h"
2  #include "stdlib.h"
3  #include "time.h"
4
5  #define MAX_STRIDE 20
6
7  int main()
8  {
9      int N = 1000000;
10     double *a;
11     a = malloc(N * MAX_STRIDE * sizeof(double));
12     double sum, rate, msec, start, end;
13
14     for (int i = 0; i < N * MAX_STRIDE; i++)
15         a[i] = 1.;
16
17     printf("stride\tsum\ttime(msec)\trate(MB/s)\n");
18
19     for (int i_stride = 1; i_stride <= MAX_STRIDE; i_stride++)
20     {
21         sum = 0.0;
22         start = (double)clock() / CLOCKS_PER_SEC;
23
24         for (int i = 0; i < N * i_stride; i += i_stride)
25             sum += a[i];
26
27         end = (double)clock() / CLOCKS_PER_SEC;
28         msec = (end - start) * 1000.0; // Time in milliseconds
29         rate = sizeof(double) * N * (1000.0 / msec) / (1024 * 1024);
30
31         printf("%d\t%f\t%f\t%f\n", i_stride, sum, msec, rate);
32     }
33     free(a);
34 }
```

Compilation

- Compile the program with O0 (without any optimization) :

```
1 gcc -O0 -o stride stride.c
```

- Compile the program with O2 (with level 2 optimization) :

```
1 gcc -O2 -o stride stride.c
```

Loop Optimizations

- Loop unrolling (partially) :

```
1 for (int i = 0; i < N; i++) {
2     sum += arr[i];
3 }
```

After unrolling :

```
1 for (int i = 0; i < N; i += 4) {
2     sum += arr[i] + arr[i + 1] + arr[i + 2] + arr[i + 3];
3 }
```

- Instruction scheduling : Reordering instructions to improve pipeline efficiency.

```
1 MUL R1, R2, R3 ; Multiply (long latency)
2 ADD R4, R1, R5 ; Add (depends on R1)
3 SUB R6, R7, R8 ; Independent subtraction
```

After reordering :

```
1 MUL R1, R2, R3 ; Multiply (long latency)
2 SUB R6, R7, R8 ; Independent subtraction (executed while MUL is
   ↪ running)
3 ADD R4, R1, R5 ; Add (by now, R1 is ready)
```

Execution and Analysis

- Run the code using -O0 and -O2 for different strides.
- Compare the execution times and bandwidths.
- Discuss the impact of loop unrolling.

Exercise 2

- Write `mxm.c` to implement the standard matrix multiplication using three nested loops.

```
1 for (int i = 0; i < n; i++)
2     for (int j = 0; j < n; j++)
3         for (int k = 0; k < n; k++)
4             c[i][j] += a[i][k] * b[k][j];
```

- Modify the loop order (jk) to optimize cache usage and improve performance.
- Compute the execution time and memory bandwidth for both versions and compare the results.
- Explain the output.

Exercise 3

- Write `mxm_bloc.c` for block matrix multiplication.
- Compute the CPU time and memory bandwidth for different block sizes.
- Determine the optimal block size. Explain why it is the best choice.

Compilation

```
gcc -O2 -o mxm_block mxm_bloc.c
```

Execution and Analysis

- Run the program with different block sizes.
- Compare the CPU time and bandwidth for each block size.
- Identify the optimal block size and justify why it provides the best performance.

Instructions

- Modify the standard matrix multiplication algorithm to process submatrices (blocks) instead of individual elements.
- Use three nested loops, but ensure that matrix elements are accessed in blocks of size $B \times B$.
- Follow this structure for blocking :
 - Divide matrices A, B, and C into blocks of size $B \times B$.
 - Compute the result for each block before moving to the next.

Exercise 4 : Memory Management and Debugging with Valgrind

- Analyze the following C program, which allocates, initializes, prints, and duplicates an array.

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4
5  #define SIZE 5
6
7  int* allocate_array(int size) {
8      int *arr = (int*)malloc(size * sizeof(int));
9      if (!arr) {
10         fprintf(stderr, "Memory_allocation_failed\n");
11         exit(EXIT_FAILURE);
12     }
13     return arr;
14 }
15
16 void initialize_array(int *arr, int size) {
17     if (!arr) return;
18     for (int i = 0; i < size; i++) {
19         arr[i] = i * 10;
20     }
21 }
22
23 void print_array(int *arr, int size) {
24     if (!arr) return;

```

```
25 printf("Array elements:");
26 for (int i = 0; i < size; i++) {
27     printf("%d", arr[i]);
28 }
29 printf("\n");
30 }
31
32 int* duplicate_array(int *arr, int size) {
33     if (!arr) return NULL;
34     int *copy = (int*)malloc(size * sizeof(int));
35     if (!copy) {
36         fprintf(stderr, "Memory allocation failed\n");
37         exit(EXIT_FAILURE);
38     }
39     memcpy(copy, arr, size * sizeof(int));
40     return copy;
41 }
42
43 void free_memory(int *arr) {
44     // add free memory line to fix the memory leak
45 }
46
47 int main() {
48     int *array = allocate_array(SIZE);
49     initialize_array(array, SIZE);
50     print_array(array, SIZE);
51     int *array_copy = duplicate_array(array, SIZE);
52     print_array(array_copy, SIZE);
53     free_memory(array);
54     return 0; // Memory leak on purpose
55 }
```

Compilation and Execution

- Compile the program with debugging symbols :

```
gcc -g -o memory_debug memory_debug.c
```

- Run the program using Valgrind to check for memory leaks :

```
valgrind --leak-check=full --track-origins=yes ./memory_debug
```

- Use Valgrind's Memcheck tool to detect memory leaks.
- Modify the program to fix memory leaks and re-run Valgrind to verify.