

# TP3 - OpenMP (Parallel Sections, Single, Master, Synchronization)

Imad Kissami

June, 2025

## Exercise 1 : Work Distribution with Parallel Sections

- Write a program that initializes an array of size  $N$  with random values.
- Use `#pragma omp sections` to divide the work :
  - Section 1 : Compute the sum of all elements.
  - Section 2 : Compute the maximum value.
  - Section 3 : Compute the standard deviation.
- Ensure that all computations run in parallel.

## Exercise 2 : Ordered Execution with Single

- Implement an OpenMP program where multiple threads generate numbers in parallel.
- Only one thread at a time should print a number using `single`.
- The numbers must be printed in ascending order.

**Example Output (using 4 threads) :**

```
1 Thread 3 generated value: 12
2 Thread 1 generated value: 27
3 Thread 2 generated value: 34
4 Thread 0 generated value: 89
```

## Exercise 3 : Exclusive Execution - Master vs Single

- Write a program where :
  - A master thread initializes a matrix.
  - A single thread prints the matrix.
  - All threads compute the sum of all elements in parallel.
- Compare execution time with and without OpenMP.

## Exercise 4 : Barrier Synchronization

- Implement a program where multiple threads execute different stages of computation :
  - Stage 1 : Read input data (only one thread should do this).
  - Stage 2 : All threads process the data in parallel.
  - Stage 3 : A single thread writes the final result.
- Use `barrier` to enforce correct execution order.

## Exercise 5 : Load Balancing with Parallel Sections

- Implement a task scheduling mechanism using parallel sections.
- Simulate three different workloads :
  - Task A (light computation)
  - Task B (moderate computation)
  - Task C (heavy computation)
- Measure the execution time and optimize the workload distribution.

## Exercise 6 : Critical vs Atomic for Shared Counters

- Implement a counter that multiple threads increment simultaneously.
- Test the program using both :
  - `critical` section
  - `atomic` directive
- Compare performance and explain the differences.

## Exercise 7 : Producer-Consumer Problem

- A producer generates values and places them in a shared buffer.
- A consumer retrieves these values and processes them.
- The two entities must be synchronized to avoid race conditions.

Sample code :

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <omp.h>
4
5  #define N 1000000
6
7  void fill_rand(int n, double *A) {
8      for (int i = 0; i < n; i++)
9          A[i] = rand() % 100;
10 }
11
12 double Sum_array(int n, double *A) {
13     double sum = 0.0;
14     for (int i = 0; i < n; i++)
15         sum += A[i];
16     return sum;
17 }
18
19 int main() {
20     double *A, sum, runtime;
21     int flag = 0; // Synchronization flag
22

```

```
23  A = (double *)malloc(N * sizeof(double));
24
25  runtime = omp_get_wtime();
26
27  fill_rand(N, A); // Producer fills the array
28
29  sum = Sum_array(N, A); // Consumer computes the sum
30
31  runtime = omp_get_wtime() - runtime;
32
33  printf("In %lf seconds, the sum is %lf\n", runtime, sum);
34
35  free(A);
36  return 0;
37 }
```