

Imad Kissami¹

¹Mohammed VI Polytechnic University, Benguerir, Morocco

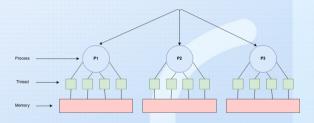
October 30, 2025



Outline of this lecture

- Distributed Memory Architectures & MPI
- Point-to-point communications
- Collective communications
- Communication modes
- Derivatives types
- Communicators

Distributed Memory Multiprocessors



- Each processor has a local memory
- Processors must communicate to access non-local data
- Parallel applications must be partitioned across
 - Processors: execution units
 - Memory: data partitioning

Message passing programming model

In the message-passing model:

- The program is written in a conventional language (e.g., C, Fortran).
- Variables are private and reside in each process's local memory.
- Communication between processes occurs via explicit message-passing subroutine calls.

Message passing concepts

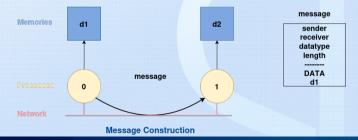
If a message is sent to a process, the process must receive it.



Message content

A message comprises:

- Data chunks passing from the sending process to the receiving process/pocesses (e.g., scalar variables, arrays).
- Metadata, including:
 - ID of the sending process, Data type and length, ID of the receiving process.

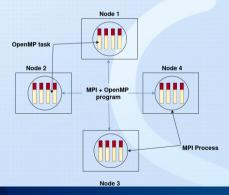


Environment

- The exchanged messages are interpreted and managed by an environment comparable to telephony, e-mail, postal mail, etc.
- The message is sent to a specified address.
- The receiving process must be able to classify and interpret the messages which are sent to it.
- The environment in question is MPI (Message Passing Interface). An MPI application is a group of autonomous processes, each executing its own code and comcommunicating via calls to MPI library subroutines.

MPI vs. OpenMP

- MPI uses a distributed memory model.
- OpenMP uses a shared memory model.



History

- MPI 1.0 (June 1994): Initial definition.
- MPI 2.0 (1997): Added features like parallel I/O.
- MPI 3.0 (2012): Introduced nonblocking collective communication.
- MPI 4.0 (2021): Added large count, partitioned communication.
- MPI 4.1 (2023): Clarifications and small extensions to MPI 4.0, including improved partitioned communication semantics, new predefined reduction operations, and clarifications for persistent collective operations.
- MPI 5.0 (June 2025): Major update introducing a standard Application Binary
 Interface (ABI) for interoperability among MPI implementations, expanded language
 bindings (including modern C++ support), improved asynchronous and persistent
 communication models, new tools interfaces, and further refinements to partitioned
 communication and large-count operations.

MPI_Init & MPI_Finalize

- Every program unit calling MPI routines has to include the header file mpi.h.
- Initialize MPI environment: MPI_Init()

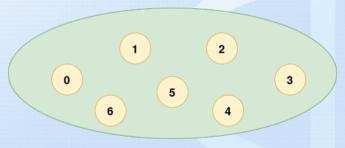
```
int MPI_Init(int *argc, char ***argv);
```

• Finalize MPI environment: MPI_Finalize()

```
int MPI_Finalize(void)
```

Communicators

- All the MPI operations occur in a defined set of processes, called communicator.
- The default communicator is MPI_COMM_WORLD, which includes all the active processes.



MPI_COMM_WORLD Communicator

Termination of a program

Sometimes, a program encounters some issue during its execution and has to stop prematurely. For example, we want the execution to stop if one of the processes cannot allocate the memory needed for its calculation. In this case, we call the MPI_Abort() subroutine instead of the Fortran instruction stop (Or exit in C).

```
int MPI_Abort(MPI_Comm comm, int error)
```

- comm: the communicator of which all the processes will be stopped; it is advised to use MPI_COMM_WORLD in general.
- error: the error number returned to the UNIX environment

The MPI library

Rank and size

 One can access to the number of processes managed by a given communicator using the MPI_Comm_size() function

```
int MPI_Comm_size(MPI_Comm comm,int *nb_procs)
```

You can get the rank of a process, within a communicator, by calling MPI_Comm_rank() function

```
int MPI_Comm_rank(MPI_Comm comm,int *rank)
```

The MPI library

Rank and size Example

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, size;

MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

printf("I am process %d among %d\n", rank, size);

MPI_Finalize();
    return 0;
}
```

```
$ mpicc -o who_am_I who_am_I.c

$ mpirun -n 2 who_am_I
I am the process 0 among 2
I am the process 1 among 2
```

Blocking Send MPI_SEND

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest,int tag, MPI_Comm comm)
```

- Sending, from the address buf, a message of count elements of type datatype, tagged tag, to the process of rank dest in the communicator comm.
- the execution remains blocked until the message can be re-written without risk of overwriting the value to be sent. In other words, the execution is blocked as long as the message has not been received.

Blocking Receive MPI_RECV

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag,MPI_Comm comm, \hookleftarrow MPI_Status *status_msg)
```

- Receiving, at the address buf, a message of count elements of type datatype, tagged tag, from the process of rank source in the communicator comm.
- status_msg stores the state of a receive operation : source, tag, code, ...
- An MPI_RECV can only be associated to an MPI_SEND if these two calls have the same envelope (source, dest, tag, comm).
- the execution remains blocked until the message content corresponds to the received message.

Blocking Send / Receive Full example

```
#include <mpi.h>
#include <stdio h>
int main(int argc, char *argv[]) {
   int rank, value, tag = 100;
   MPI_Status status;
   MPI Init(&argc, &argv); // Initialize MPI environment
   MPI Comm rank(MPI COMM WORLD, &rank); // Get process rank
   if (rank == 2) {
      value = 1000:
      MPI Send(&value, 1, MPI INT, 5, tag, MPI COMM WORLD):
   } else if (rank == 5) {
      MPI Recv(&value, 1, MPI INT, 2, tag, MPI COMM WORLD, &status):
      printf("I, process 5, received %d from process 2.\n", value);
   MPI Finalize(): // Finalize MPI environment
   return 0:
```

```
mpirun -n 6 point_to_point
I, process 5, I received 1000 from process 2.
```

C MPI Datatypes

MPI Datatype	C Datatype
MPI_CHAR	signed char
MPI_SHORT	signed short
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	Raw byte data

Table: Mapping of MPI datatypes to C datatypes

Other possibilities

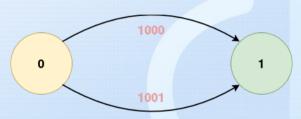
- Wildcard Parameters
 - The sender process rank can be replaced with MPI_ANY_SOURCE.
 - The message tag can be replaced with MPI_ANY_TAG.
- Dummy Processes
 - Communications with the process rank MPI_PROC_NULL have no effect.
- Ignoring Status
 - Use the predefined constant MPI_STATUS_IGNORE instead of the status variable if the status information is not needed.
- Derived Datatypes
 - You can send more complex data structures by creating derived datatypes.
- Simultaneous Send and Receive
 - Use operations like MPI_Sendrecv() or MPI_Sendrecv_replace() for simultaneous send and receive actions.

Simultaneous send and receive MPI_SENDRECV

```
int MPI_Sendrecv(const void *sendbuf, int sendcount, MPI_Datatype sendtype, int dest, int sendtag←
    , void *recvbuf, int recvcount, MPI_Datatype recvtype, int source, int recvtag, MPI_Comm ←
    comm, MPI_Status *status_msg)
```

- Sending, from the address sendbuf, a message of sendcount elements of type sendtype, tagged sendtag, to the process dest in the communicator comm;
- Receiving, at the address recvbuf, a message of recvcount elements of type recvtype, tagged recvtag, from the process source in the communicator comm.
- Here, the receiving zone recybuf must be different from the sending zone sendbuf.

Simultaneous send and receive MPI_SENDRECV



sendrecv Communication between the Processes 0 and 1

Simultaneous send and receive MPI_SENDRECV: Full example

```
#include <mpi.h>
#include <stdio.h>

int main(int argc,char *argv[]) {
    int rank,value, num_proc, message, tag=110;
    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    num_proc=(rank+1)%2;
    message = rank+1000;
    MPI_Sendrecv(&message,1,MPI_INT,num_proc,tag,&value,1,MPI_INT,
    num_proc,tag,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
    printf("I, process %d, I received %d from process %d.\n", rank,value,num_proc);
    MPI_Finalize();
}
```

```
mpirun -n 2 simultaneoussendrecv
I, process 0, I received 1001 from the process 1.
I, process 1, I received 1000 from the process 0.
```

Simultaneous send and receive MPI_SENDRECV: Remarks

In the case of a synchronous implementation of the MPI_SEND() subroutine, if we replace the MPI_SENDECV() subroutine in the example above by MPI_SEND() followed by MPI_RECV(), the code will deadlock. Indeed, each of the two processes will wait for a receipt confirmation, which will never come because the two sending operations would stay suspended.

```
val = rank+1000;
MPI_Send(&val,1,MPI_INT,num_proc,tag,MPI_COMM_WORLD);
MPI_Recv(value,1,MPI_INT,num_proc,tag,MPI_COMM_WORLD,&statut);
```

Simultaneous send and receive MPI_SENDRECV_REPLACE

```
int MPI_Sendrecv_replace(void *buf, int count, MPI_Datatype datatype, int dest, int sendtag, int \hookleftarrow source, int recvtag, MPI_Comm comm, MPI_Status *status)
```

- Sending, from the address buf, a message of count elements of type datatype, tagged sendtag, to the process dest in the communicator comm;
- Receiving a message at the same address, with same count elements and same datatype, tagged recvtag, from the process source in the communicator comm.
- Contrary to the usage of MPI_SENDRECV, the receiving zone is the same here as the sending zone buf.

Simultaneous send and receive MPI_SENDRECV_REPLACE : full example

```
int main(int argc, char *argv[]) {
   int rank, tag = 11, m = 2;
   int A[m][m]:
   MPI Status status:
   MPI Init(&argc, &argv);
   MPI Comm rank(MPI COMM WORLD, &rank):
   if (rank == 0) {
      A[0][0] = 1; A[0][1] = 2; A[1][0] = 3; A[1][1] = 4;
      MPI Send(A. 2. MPI INT. 1. tag. MPI COMM WORLD):
   } else {
      MPI_Recv(&(A[0][1]), 2, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
      MPI COMM WORLD, &status):
      printf("I process %d, received 2 elements from process %d with tag %d. "
      "The elements are %d %d.\n", rank, status.MPI SOURCE, status.MPI TAG, A[0][1], A[1][0]):
   MPI_Finalize();
```

Simultaneous send and receive MPI_SENDRECV_REPLACE : full example

```
mpirun -n 2 ./simultaneoussendrecv_replace
I process 1, received 2 elements from process 0 with tag 11. The elements are 1 2.
```

General concepts

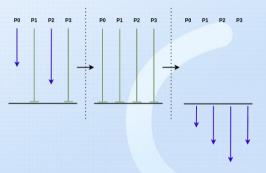
- Collective communications allow making a series of point-to-point communications in one single call.
- A collective communication always concerns all the processes of the indicated communicator.
- For each process, the call ends when its participation in the collective call is completed, in the sense of point-to-point communications (therefore, when the concerned memory area can be changed).
- The management of tags in these communications is transparent and system-dependent. Therefore, they are never explicitly defined during calls to subroutines. An advantage of this is that collective communications never interfere with point-to-point communications.

Types of collective communications

- One which ensures global synchronizations: MPI_Barrier()
- Ones which only transfer data:
 - Global distribution of data: MPI_Bcast()
 - Selective distribution of data: MPI_Scatter()
 - Collection of distributed data: MPI_Gather()
 - Collection of distributed data by all the processes: MPI_Allgather()
 - Collection and selective distribution by all the processes of distributed data: MPI_Alltoall()
- Ones which, in addition to the communications management, carry out operations on the transferred data:
 - Reduction operations (sum, product, maximum, minimum, etc.), whether of a predefined or personal type: MPI_Reduce()
 - Reduction operations with distributing of the result (this is in fact equivalent to an MPI_Reduce() followed by an MPI_Bcast()): MPI_Allreduce()



Global synchronization: MPI_Barrier()



Global Synchronization : MPI_BARRIER()

int MPI_Barrier(MPI_Comm comm)

Global distribution: MPI_BCAST()

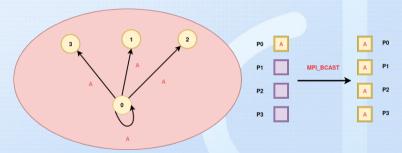


Figure: Global distribution: MPI_Bcast()

Global distribution: MPI_Bcast()

```
int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm)
```

- Send, starting at position buffer, a message of count element of type datatype, by the root process, to all the members of communicator comm.
- Receive this message at position buffer for all the processes other than the root.

Global distribution: MPI_Bcast() Full example 1

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, value;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 2) value = rank+1000

    MPI_Bcast(&value,1,MPI_INT,2,MPI_COMM_WORLD);
    printf("I, process %d, received %d of process 2\n", rank,value);

    MPI_Finalize();
}
```

```
mpirun -n 3 bcast

I, process 2, received 1002 of process 2

I, process 1, received 1002 of process 2

I, process 0, received 1002 of process 2
```

Selective distribution: MPI_Scatter()

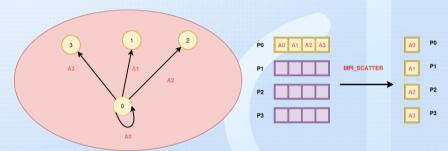


Figure: Selected distribution: MPI_Scatter()

Selective distribution: MPI_Scatter()

```
int MPI_Scatter(const void *sendbuf,int sendcount, MPI_Datatype sendtype,void *recvbuf, int ←
    recvcount,MPI_Datatype recvtype, int root,MPI_Comm comm)
```

- Scatter by process root, starting at position sendbuf, message sendcount element of type sendtype, to all the processes of communicator comm.
- Receive this message at position recvbuf, of recvcount element of type recvtype for all processes of communicator comm.
- The couples (sendcount, sendtype) and (recvcount, recvtype) must represent the same quantity of data.
- Data are scattered in chunks of same size; a chunk consists of sendcount elements of type sendtype.
- The i-th chunk is sent to the i-th process.

Selective distribution: MPI_SCATTER() Full example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib h>
int main(int argc, char *argv[]) {
   int nb values = 8, rank, nb procs, block length, i:
   float *values. *recydata:
   MPI Init(&argc, &argv):
   MPI Comm size(MPI COMM WORLD, &nb procs);
   MPI Comm rank(MPI COMM WORLD, &rank);
   block_length = nb_values / nb_procs;
   recvdata = (float *)malloc(block length * sizeof(float));
   if (rank == 2) {
      values = (float *)malloc(nb values * sizeof(float));
      for (i = 0; i < nb \ values; i++) \ values[i] = 1001.0 + i;
      printf("Process %d sends the values array: ", rank);
      for (i = 0: i < nb values: i++) printf("%f ", values[i]):
      printf("\n"):
```

Selective distribution: MPI_SCATTER() Full example

```
mpirun -n 4 scatter
I, process 2 send the values array : 1001. 1002. 1003. 1004. 1005. 1006. 1007. 1008.
I, process 0, received 1001. 1002. from processus 2
I, process 1, received 1003. 1004. from processus 2
I, process 3, received 1007. 1008. from processus 2
I, process 2, received 1005. 1006. from processus 2
I, process 2, received 1005. 1006. from processus 2
```

Selective collection: MPI_Gather()

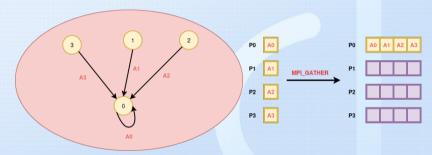


Figure: Collection : MPI_Gather()

Selective collection: MPI_Gather()

```
int MPI_Gather(const void*sendbuf,int sendcount, MPI_Datatype sendtype,void *recvbuf, int ←
    recvcount,MPI_Datatype recvtype, int root,MPI_Comm comm)
```

- Send for each process of communicator comm, a message starting at position sendbuf, of sendcount element type sendtype
- Collect all these messages by the root process at position recvbuf, recvcount element of type recvtype.
- The couples (sendcount, sendtype) and (recvcount, recvtype) must represent the same size of data.
- The data are collected in the order of the process ranks.

Selective collection: MPI_GATHER() Full example

```
//gather.c
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]) {
   int nb values=6.rank.nb procs.block length.i:
   float recvdata[nb_values],*values;
   MPI Init(&argc, &argv):
   MPI Comm size(MPI COMM WORLD, &nb procs);
   MPI Comm rank(MPI COMM WORLD, &rank):
   block length = nb values/nb procs:
   values = (float *)malloc(block length * sizeof(float)):
   for (i = 0; i < block_length; i++) values[i]=1001.+rank*block_length+i;</pre>
   printf("I, process %d sent my values array : ",rank);
```

Selective collection: MPI_GATHER() Full example

Global collection: MPI_Allgather()

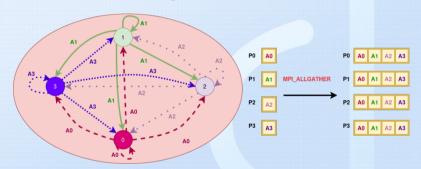


Figure: Gather-to-all: MPI_Allgather()

Global collection: MPI_Allgather()

```
int MPI_Allgather(const void *sendbuf,int sendcount, MPI_Datatype sendtype,void *recvbuf, int \hookleftarrow recvcount,MPI_Datatype recvtype, MPI_Comm comm)
```

- Corresponds to an MPI_Gather() followed by an MPI_Bcast()
- Send by each process of communicator comm, a message starting at position sendbuf, of sendcount element, type sendtype.
- Collect all these messages, by all the processes, at position recvbuf of recvcount element type recvtype.
- The couples (sendcount, sendtype) and (recvcount, recvtype) must represent the same data size.
- The data are gathered in the order of the process ranks.

Global collection: MPI_Allgather() Full example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib h>
int main(int argc, char *argv[]) {
   int nb values = 6, rank, nb procs, block length, i;
   float recydata[nb values]. *values:
   MPI Init(&argc, &argv):
   MPI_Comm_size(MPI_COMM_WORLD, &nb_procs);
   MPI_Comm_rank(MPI_COMM_WORLD, &rank);
   block length = nb values / nb procs:
   values = (float *)malloc(block length * sizeof(float));
   for (i = 0; i < block_length; i++) values[i] = 1001. + rank * block_length + i;</pre>
   MPI Allgather(values, block length, MPI FLOAT, recvdata, block length, MPI FLOAT, ←
         MPI COMM WORLD):
  /* print output */
   MPI Finalize():
```

Global collection: MPI_Allgather() Full example

```
mpirun -n 3 allgather

I, process 0, received 1001.000000 1002.000000 1003.000000 1004.000000 1005.000000 1006.000000
I, process 1, received 1001.000000 1002.000000 1003.000000 1004.000000 1005.000000 1006.000000
I, process 2, received 1001.000000 1002.000000 1003.000000 1004.000000 1005.000000 1006.000000
```

Selective collection: MPI_Gatherv()

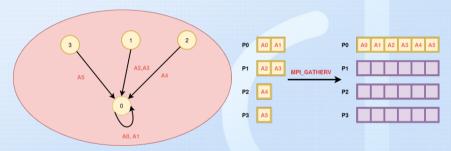


Figure: Extended gather: MPI_Gatherv()

Selective collection: MPI_Gatherv()

```
int MPI_Gatherv(const void *sendbuf,int sendcount,MPI_Datatype sendtype, void *recvbuf,const int *
    recvcounts,const int *displs, MPI_Datatype recvtype,root,MPI_Comm comm)
```

- This is an MPI_Gather() where the size of mess. can be different among procs
- The i-th proc of the communicator comm sends to root, a message starting at position sendbuf, of sendcount element of type sendtype, and receives at position recvbuf, of recvcounts(i) element of type recvtype, with a displacement of displs(i)
- The couples (sendcount, sendtype) of the i-th process and (recvcounts(i), recvtype) of process root must be such that the data size sent and received is the same.

Selective collection: MPI_Gatherv() Full example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib h>
int main(int argc, char *argv[]) {
   int nb values = 10, rank, nb procs, block length, remainder, i:
   float recydata[nb values]:
   float *values:
   int *nb elements received. *displacement:
   MPI_Init(&argc, &argv);
   MPI Comm size(MPI COMM WORLD, &nb procs);
   MPI Comm rank(MPI COMM WORLD, &rank);
   block length = nb values / nb procs;
   remainder = nb values % nb procs:
   if (rank < remainder) block length = block length + 1:</pre>
   values = (float *)malloc(block_length * sizeof(float));
   for (i = 0: i < block length: i++)
   values[i] = 1001. + rank * (nb values / nb procs) + (rank < remainder ? rank : remainder) + i;</pre>
```

Selective collection: MPI_Gatherv() Full example

```
printf("I, process %d send my values array: ", rank);
for (i = 0: i < block length: i++) {
   printf("%f ", values[i]):
printf("\n");
if (rank == 2) {
   nb elements received = (int *)malloc(nb procs * sizeof(int)):
   displacement = (int *)malloc(nb procs * sizeof(int));
   nb elements received[0] = nb values / nb procs;
   if (remainder > 0) nb elements received[0] = nb elements received[0] + 1;
   displacement[0] = 0;
   for (i = 1; i < nb_procs: i++) {
      displacement[i] = displacement[i - 1] + nb elements received[i - 1];
      nb elements received[i] = nb values / nb procs;
      if (i < remainder) nb elements received[i] = nb elements received[i] + 1:
MPI Finalize():
```

Selective collection: MPI_Gatherv() Full example

```
mpirun -n 4 gatherv

I, process 2 send my values array: 1007.000000 1008.000000

I, process 3 send my values array: 1009.000000 1010.000000

I, process 1 send my values array: 1004.000000 1005.000000 1006.000000

I, process 0 send my values array: 1001.000000 1002.000000 1003.000000
```

Global collection & distribution: MPI_Alltoall()

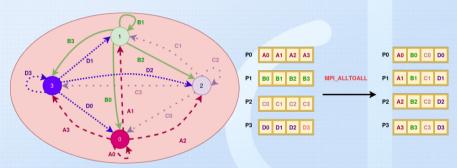


Figure: Collection and distribution : : MPI_Alltoall()

Global collection & distribution: MPI_Alltoall()

```
\begin{tabular}{ll} \textbf{int MPI\_Alltoall(const void *sendbuf,int sendcount, MPI\_Datatype sendtype,void *recvbuf, int} & recvcount, MPI\_Datatype recvtype, MPI\_Comm comm) \\ \end{tabular}
```

- the i-th process sends its j-th chunk to the j-th process which places it in its i-th chunk.
- The couples (sendcount, sendtype) and (recvcount, recvtype) must be such that they represent equal data sizes.

Global collection & distribution: MPI_Alltoall() Full example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib h>
int main(int argc, char *argv[]) {
   int nb_values = 8;
   int rank, nb_procs, block_length, i;
   int recvdata[nb values], values[nb values];
   MPI Init(&argc, &argv):
   MPI Comm size(MPI COMM WORLD, &nb procs):
   MPI Comm rank(MPI COMM WORLD, &rank);
   for (i = 0: i < nb values: i++)
   values[i] = 1001. + rank * nb values + i;
   block length = nb values / nb procs:
   printf("I, process %d sent mv values array: ", rank);
   for (i = 0; i < nb values; i++) printf("%d ", values[i]);</pre>
   printf("\n"):
```

Global collection & distribution: MPI_Alltoall() Full example

```
MPI_Alltoall(values, block_length, MPI_INT, recvdata, block_length, MPI_INT, MPI_COMM_WORLD);
printf("I, process %d, received: ", rank);
for (i = 0; i < nb_values; i++) printf("%d ", recvdata[i]);
printf("\n");
MPI_Finalize();
return 0;
}</pre>
```

```
mpirun -n 4 alltoall

I, process 0 sent my values array: 1001 1002 1003 1004 1005 1006 1007 1008

I, process 0, received: 1001 1002 1009 1010 1017 1018 1025 1026

I, process 1 sent my values array: 1009 1010 1011 1012 1013 1014 1015 1016

I, process 1, received: 1003 1004 1011 1012 1019 1020 1027 1028

I, process 2 sent my values array: 1017 1018 1019 1020 1021 1022 1023 1024

I, process 2, received: 1005 1006 1013 1014 1021 1022 1029 1030

I, process 3 sent my values array: 1025 1026 1027 1028 1029 1030 1031 1032

I, process 3, received: 1007 1008 1015 1016 1023 1024 1031 1032
```

Global Reduction

- A reduction is an operation applied to a set of elements in order to obtain one single value. Typical examples are the sum of the elements of a vector (SUM(A(:))) or the search for the maximum value element in a vector (MAX(V(:))).
- MPI proposes high-level subroutines in order to operate reductions on data
 distributed on a group of processes. The result is obtained on only one process (
 MPI_Reduce()) or on all the processes (MPI_Allreduce(), which is in fact equivalent to
 an MPI_Reduce() followed by an MPI_Bcast()).
- If several elements are implied by process, the reduction function is applied to each one of them (for instance to each element of a vector).

Global Reduction

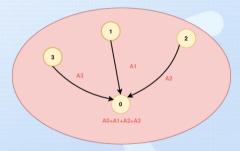


Figure: Distributed reduction (sum)

Global Reduction: Operations

	Name	Operation
ĺ	MPI_SUM	Sum of elements
	MPI_PROD	Product of elements
	MPI_MAX	Maximum of elements
	MPI_MIN	Minimum of elements
	MPI_MAXLOC	Maximum of elements and location
	MPI_MINLOC	Minimum of elements and location
	MPI_LAND	Logical AND
	MPI_LOR	Logical OR
	MPI_LXOR	Logical exclusive OR
		_

Table: Global Reduction available operations

Global Reduction: MPI_Reduce()

```
int MPI_Reduce(const void*sendbuf,void *recvbuf,int count, MPI_Datatype datatype,MPI_Op op,int ←
    root,
MPI_Comm comm)
```

- Distributed reduction of count elements of type datatype, starting at position sendbuf, with the operation op from each process of the communicator comm,
- Return the result at position recybuf in the process root

Global Reduction: MPI_Reduce() Full example

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
   int rank, nb_procs, value, sum;
   MPI_Init(&argc, &argv);
   MPI Comm size(MPI COMM WORLD, &nb procs);
   MPI Comm rank(MPI COMM WORLD, &rank);
   if (rank == 0)
   value = 1000:
   else
   value = rank:
   MPI Reduce(&value, &sum, 1, MPI INT, MPI SUM, 0, MPI COMM WORLD):
   if (rank == 0)
   printf("I, process 0, have the global sum value %d\n", sum);
   MPI_Finalize();
```

Global Reduction: MPI_Reduce() Full example

```
mpirun -n 8 reduce
```

I, process 0, have the global sum value 1028

Global Reduction: MPI_Allreduce()

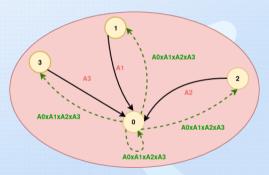


Figure: Distributed reduction (product) with distribution of the result

Global Reduction: MPI_Allreduce()

- Distributed reduction of count elements of type datatype starting at position sendbuf, with the operation op from each process of the communicator comm,
- Write the result at position recvbuf for all the processes of the communicator comm.

Global Reduction: MPI_Allreduce() Full example

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
   int rank, nb procs, value, product;
   MPI Init(&argc, &argv);
   MPI_Comm_size(MPI_COMM_WORLD, &nb_procs);
   MPI Comm rank(MPI COMM WORLD, &rank);
   if (rank == 0)
   value = 10;
   else
   value = rank:
   MPI Allreduce(&value, &product, 1, MPI INT, MPI PROD, MPI COMM WORLD);
   printf("I, process %d, received the value of the global product %d\n".
   rank, product):
   MPI Finalize():
```

Global Reduction: MPI_Allreduce() Full example

```
mpirun -n 8 allreduce

I, process 7, received the value of the global product 50400
I, process 0, received the value of the global product 50400
I, process 1, received the value of the global product 50400
I, process 2, received the value of the global product 50400
I, process 3, received the value of the global product 50400
I, process 4, received the value of the global product 50400
I, process 5, received the value of the global product 50400
I, process 6, received the value of the global product 50400
I, process 6, received the value of the global product 50400
```

Additions

- MPI_Scan() allows making partial reductions by considering, for each process, the previous processes of the communicator and itself.
- MPI_Exscan() is the exclusive version of MPI_Scan(), which is inclusive.
- The MPI_Op_create() and MPI_Op_free() subroutines allow personal reduction operations.
- For each reduction operation, the keyword MPI_IN_PLACE can be used to keep the
 result in the same place as the sending buffer (but only for the rank(s) that will
 receive results). Example:

```
MPI_Allreduce(MPI_IN_PLACE, sendrecvbuf,...);
```

Additions

- Similarly to what we have seen for MPI_Gatherv() with respect to MPI_Gather(), the
 following subroutines extend their respective counterparts to handle cases where
 processes have different numbers of elements to transmit or gather:
 - MPI_Scatterv() extends MPI_Scatter().
 - MPI_Allgatherv() extends MPI_Allgather().
 - MPI_Alltoallv() extends MPI_Alltoall().
- MPI_Alltoallw() is the version of MPI_Alltoallv() that enables dealing with heterogeneous elements. It does so by expressing the displacements in bytes instead of elements.

Point-to-Point Send Modes

Mode	Blocking	Non-blocking
Standard Send	MPI_Send()	MPI_Isend()
Synchronous Send	MPI_Ssend()	<pre>MPI_Issend()</pre>
Buffered send	MPI_Bsend()	<pre>MPI_Ibsend()</pre>
Receive	MPI_Recv()	MPI_Irecv()

Table: Point-to-Point Send Modes

General concepts

Blocking call

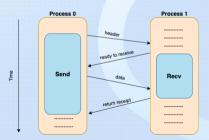
- A call is blocking if the memory space used for the communication can be reused immediately after the exit of the call
- The data sent can be modified after the call.
- The data received can be read after the call.

Synchronous sends

A synchronous send involves a synchronization between the involved processes. A send cannot start until its receive is posted. There can be no communication before the two processes are ready to communicate.

Rendezvous Protocol

The rendezvous protocol is generally the protocol used for synchronous sends (implementation-dependent). The return receipt is optional.



Interface of : MPI_SSEND()

```
int MPI_Ssend(const void* values,int count, MPI_Datatype msgtype, int dest,int tag, MPI_Comm comm←
)
```

- Advantages of synchronous mode
 - Low resource consumption (no buffer)
 - Rapid if the receiver is ready (no copying in a buffer)
 - Knowledge of receipt through synchronization
- Disadvantages of synchronous mode
 - Waiting time if the receiver is not there/not ready
 - · Risk of deadlocks

Deadlock example

In this example, there is a deadlock because we are in synchronous mode. The two processes are blocked on the MPI_Ssend() call because they are waiting for the MPI_RECV() of the other process. However, the MPI_Recv() call can only be made after the unblocking of the MPI_Ssend() call.

```
int main(int argc, char *argv[]) {
   int rank, num_proc, tmp, value;
   int tag = 110;

MPI_Init(&argc, &argv);
   MPI_Comm_rank(MPI_COMM_WORLD, &rank);

/* On suppose qu'il y a exactement 2 processus */
   num_proc = (rank + 1) % 2;

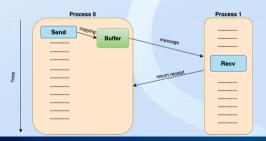
tmp = rank + 1000;
   MPI_Seend(&tmp, 1, MPI_INT, num_proc, tag, MPI_COMM_WORLD);
   MPI_Recv(&value, 1, MPI_INT, num_proc, tag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

printf("Moi, processus %d, j'ai reçu %d du processus %d\n", rank, value, num_proc);

MPI_Finalize();
}
```

Buffered sends

- A buffered send implies the copying of data into an intermediate memory space.
 There is then no coupling between the two processes of communication. Therefore, the return of this type of send does not mean that the receive has occurred.
- Protocol with user buffer on the sender side
 In this approach, the buffer is on the sender side and is managed explicitly by the
 application. A buffer managed by MPI can exist on the receiver side. Many variants
 are possible. The return receipt is optional.



Interface of: MPI_Bsend()

- The buffers have to be managed manually (with calls to MPI_Buffer_attach() and MPI_Buffer_detach()). Message header size needs to be taken into account when allocating buffers (by adding the constant MPI_Bsend_overhead() for each message occurrence).
- Interfaces

No Deadlock example

In the following example, we don't have a deadlock because we are in buffered mode. After the copy is made in the buffer, the MPI_Bsend() call returns and then the MPI_Recv() call is made.

```
int main(int argc, char *argv[]) {
   int rank, num_proc, tmp, value, bufsize, overhead, typesize;
   int tag = 110, nb_elt = 1, nb_msg = 1;
   int *buffer:
   MPI Init(&argc, &argv);
   MPI Comm rank(MPI COMM WORLD, &rank);
   MPI Type size(MPI INT, &typesize);
   /* Convert MPI BSEND OVERHEAD size (bytes) to integer number */
   overhead = (int)(1 + (MPI BSEND OVERHEAD * 1.0) / typesize);
   buffer = (int *)malloc(nb_msg * (nb_elt + overhead) * sizeof(int));
   bufsize = typesize * nb msg * (nb elt + overhead);
   MPI Buffer attach(buffer, bufsize):
   /* We assume to have exactly 2 processes */
   num_proc = (rank + 1) % 2;
   tmp = rank + 1000:
```

No Deadlock example (continuation of the code)

```
MPI_Bsend(&tmp, nb_elt, MPI_INT, num_proc, tag, MPI_COMM_WORLD);
MPI_Recv(&value, nb_elt, MPI_INT, num_proc, tag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

printf("I, process %d received %d from process %d\n", rank, value, num_proc);

MPI_Buffer_detach(&buffer, &bufsize);
MPI_Finalize();
}
```

```
mpirun -n 2 bsend

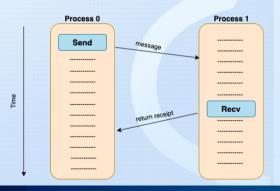
I, process 0 received 1001 from process 1
I, process 1 received 1000 from process 0
```

Interface of : MPI_Send()

- Standard sends
 A standard send is made by calling the MPI_Send() subroutine. In most implementations, the mode is buffered (eager) for small messages but is synchronous for larger messages.
- Interfaces

The eager protocol

- The eager protocol is often used for standard sends of small-size messages.
- It can also be used for sends with MPI_Bsend() for small messages (implementation-dependent) and by bypassing the user buffer on the sender side.



Number of received elements

- In MPI_Recv() call, the count argument in the standard is the number of elements in the buffer buf.
- This number must be greater than the number of elements to be received.
- When it is possible, for increased clarity, it is adviced to put the number of elements to be received.
- We can obtain the number of elements received with MPI_Get_count() and the msgstatus argument returned by the MPI_Recv() call

```
int MPI_Get_count(MPI_Status *msgstatus, MPI_Datatype msgtype,int *count)
```

Number of received elements MPI_Probe allows incoming messages to be checked for, without actually receiving them.

```
int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status)
```

A common use of MPI_Probe is to allocate space for a message before receiving it.

```
MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, comm, &status);
MPI_Get_count(&status, MPI_INT, &maggsize);
buf = (int*) malloc(msgsize*sizeof(int));
MPI_Recv(buf, msgsize, MPI_INT, status.MPI_SOURCE,
status.MPI_TAG, comm, MPI_STATUS_IGNORE);
```

The overlap of communications by computations:

- A method which allows executing communication operations in the background while the program continues to operate.
- If the hardware and software architecture allows it, it is possible to hide all or part of the communications costs.
- The computation-communication overlap can be seen as an additional level of parallelism.

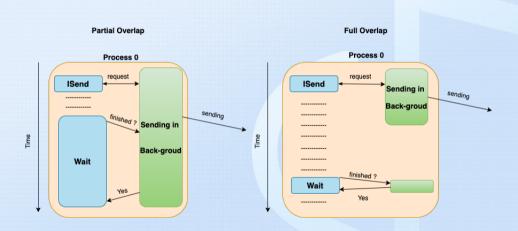
MPI Approach:

 This approach is used in MPI by using nonblocking subroutines: MPI_Isend(), MPI_Irecv() and MPI_Wait()

Nonblocking Communication:

- A nonblocking call returns very quickly but does not authorize the immediate re-use
 of the memory space which was used in the communication.
- It is necessary to make sure that the communication is fully completed (with MPI_Wait(), for example) before using it again.





Interface of: MPI_Isend(), MPI_Irecv()

Interface of : MPI_Wait()

MPI_WAIT() wait for the end of a communication, MPI_TEST() is the nonblocking version.

```
int MPI_Wait(MPI_Request *req, MPI_Status *statut)
int MPI_Test(MPI_Request *req, int *flag, MPI_Status *statut)
```

Interface of : MPI_Waitall()

MPI_Waitall() (MPI_Testall()) await the end of all communications.

```
int MPI_Waitall(int count, MPI_Request reqs[],MPI_Status statuts[])
int MPI_Testall(int count, MPI_Request reqs[],int *flag, MPI_Status statuts[])
```

Interface of : MPI_Waitall()

MPI_Waitany() wait for the end of one communication, MPI_Testany() is the nonblocking version.

```
int MPI_Waitany(int count,MPI_Request reqs[],int *indice,MPI_Status *statut)
int MPI_Testany(int count,MPI_Request reqs[],int *indice,int *flag,MPI_Status *statut)
```

MPI_Waitsome() wait for the end of at least one communication, MPI_Testsome() is the nonblocking version.

```
int MPI_Waitsome(int count,MPI_Request reqs[],int *outcount,int *indices,MPI_Status *statuses)
int MPI_Testsome(int count,MPI_Request reqs[],int *outcount,int *indices,MPI_Status *statuses)
```

Request management

- After a call to a blocking wait function (MPI_WAIT(), MPI_WAITALL(),...), the request argument is set to MPI_REQUEST_NULL.
- The same for a nonblocking wait when the flag is set to true.
- A wait call with a MPI_REQUEST_NULL request does nothing.

Contiguous datatypes: MPI_Type_contiguous()

 MPI_Type_contiguous() creates a data structure from a homogenous set of existing datatypes contiguous in memory.

```
int MPI_Type_contiguous(int count,MPI_Datatype old_type,MPI_Datatype *new_type)

MPI_Type_contiguous(6,MPI_INT,&new_type);
```

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18

MPI_Type_contiguous() subroutine

MPI_Type_commit() and MPI_Type_free()

 Before using a new derived datatype, it is necessary to validate it with the MPI_Type_commit() subroutine.

```
int MPI_Type_commit(MPI_Datatype *new_type)
```

• The freeing of a derived datatype is made by using the MPI_Type_free() subroutine.

```
int MPI_Type_free(MPI_Datatype *new_type)
```

MPI_Type_contiguous(): example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
int main(int argc,char *argv[]) {
   int rank, i, j;
   int nb lines=6,nb_columns=5, tag=100;
   float a[nb_lines][nb_columns];
   MPI Datatype type line:
   MPI Status statut:
   MPI Init(&argc,&argv);
   MPI Comm rank(MPI COMM WORLD, &rank);
   /* Initialization of the matrix on each process */
   for(i=0:i<nb lines:i++)</pre>
      for(j=0;j<nb columns;j++)</pre>
          a[i][j]=rank;
   /* Definition of the type_line datatype */
   MPI_Type_contiguous(nb_columns,MPI_FLOAT,&type_line);
   /* Validation of the type_line datatype */
   MPI_Type_commit(&type_line);
```

MPI_Type_contiguous(): example

```
/* Sending of the first line */
if (rank == 0) {
   MPI Send(a.1.type line.1.tag.MPI COMM WORLD):
} else {
   /* Receiving in the last line */
   MPI_Recv(&(a[nb_lines-1][0]),nb_columns,MPI_FLOAT,0,tag,
   MPI_COMM_WORLD,&statut); }
/* Print the matrix */
printf("Process %d matrix:\n", rank);
for (i = 0; i < nb lines; i++) {</pre>
   for (j = 0; j < nb_columns; j++) {</pre>
       printf("%.1f ", a[i][j]);
   printf("\n");
printf("\n"):
/* Free the datatype */
MPI Type free(&type line):
MPI_Finalize();
```

MPI_Type_contiguous(): example

Constant stride: MPI_Type_vector()

 MPI_Type_vector() creates a data structure from a homogenous set of existing datatypes separated by a constant stride in memory. The stride is given in number of elements.

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18

MPI_Type_vector() subroutine

MPI_Type_vector(): example

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
int main(int argc,char *argv[]) {
   int rank, i, j;
   int nb lines=6,nb columns=5, tag=100;
   float a[nb lines][nb columns];
   MPI_Datatype type_column;
   MPI Status statut:
   MPI Init(&argc,&argv);
   MPI Comm rank(MPI COMM WORLD, &rank):
   /* Initialisation of the matrix on each process */
   for(i=0:i<nb lines:i++)</pre>
      for(j=0;j<nb_columns;j++)</pre>
          a[i][i]=rank:
   /* Definition of the datatype type_column */
   MPI Type vector(nb lines.1.nb columns.MPI FLOAT.&type column);
   /* Validation of the datatype type_column */
   MPI Type commit(&type column):
```

MPI_Type_vector(): example

```
/* Sending of the first column */
if (rank == 0) {
    MPI Send(&(a[1][0]).nb lines.MPI FLOAT.1.tag.MPI COMM WORLD):
 } else {
/* Reception in the last column */
MPI_Recv(&(a[0][nb_columns-1]),1,type_column,0,tag, MPI_COMM_WORLD,&statut); }
/* Print the matrix */
printf("Process %d matrix:\n", rank);
for (i = 0; i < nb lines; i++) {
    for (j = 0: j < nb columns: j++) {</pre>
       printf("%.1f ", a[i][j]);
    printf("\n");
printf("\n");
/* Free the datatype type_column */
MPI Type free(&type colonne):
MPI_Finalize();
```

MPI_Type_vector(): example

MPI_Type_vector(): send block

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
int main(int argc,char *argv[]) {
   int rank,i,j, nb lines=6,nb columns=5, tag=100;
   int nb lines block=3,nb columns block=2;
   float a[nb lines][nb columns];
   MPI_Datatype type_block;
   MPI Status statut:
   MPI Init(&argc,&argv);
   MPI Comm rank(MPI COMM WORLD, &rank):
   /* Initialization of the matrix on each process */
   for(i=0:i<nb lines:i++)</pre>
      for(j=0;j<nb columns;j++)</pre>
          a[i][i]=rank:
   /* Creation of the datatype type block */
   MPI Type vector(nb lines block,nb columns block,nb columns, MPI FLOAT.&type block):
   /* Validation of the datatype type block */
   MPI Type commit(&type block):
```

MPI_Type_vector(): send block

```
/* Sending of a block */
if (rank == 0) {
   MPI Send(a,1,type block,1,tag,MPI COMM WORLD);
} else {
   /* Reception of a block */
   MPI_Recv(&(a[nb_lines-3][nb_columns-2]),1,type_block,0,tag,
   MPI_COMM_WORLD,&statut); }
/* Print the matrix */
printf("Process %d matrix:\n", rank);
for (i = 0; i < nb lines; i++) {
   for (j = 0; j < nb_columns; j++) {</pre>
       printf("%.1f ", a[i][i]);
   printf("\n"):
printf("\n");
/* Free the datatype type_column */
MPI_Type_free(&type_block);
MPI_Finalize();
```

MPI_Type_vector(): send block

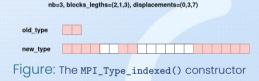
Constant stride: MPI_Type_create_hvector()

- MPI_Type_create_hvector() creates a data structure from a homogenous set of
 existing datatype separated by a constant stride in memory. The stride is given in
 bytes.
- This call is useful when the old type is no longer a base datatype (MPI_INTEGER, MPI_REAL,...) but a more complex datatype constructed by using MPI subroutines, because in this case the stride can no longer be given in number of elements.

Homogenous datatypes of variable strides

- MPI_Type_indexed(): create a data structure composed of a sequence of blocks containing a variable number of elements separated by a variable stride in memory. The stride is given in number of elements.
- MPI_Type_create_hindexed() has the same functionality as MPI_Type_indexed() except that the strides separating two data blocks are given in bytes. This subroutine is useful when the old datatype is not an MPI base datatype (MPI_INTEGER, MPI_REAL, ...). We cannot therefore give the stride in number of elements of the old datatype.
- For MPI_Type_create_hindexed(), as for MPI_Type_create_hvector(), use
 MPI_Type_size() or MPI_Type_get_extent() in order to obtain in a portable way the size of the stride in bytes.

Homogenous datatypes of variable strides: MPI_Type_indexed()



int MPI_Type_indexed(int nb,const int block_lengths[],const int displacements[], MPI_Datatype \(\to \) old_type,MPI_Datatype *new_type)

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
   int rank, size = 10;
   double data[size];
   int block_lengths[] = {2, 1, 3};
   int displacements[] = {0, 3, 7};
   MPI_Datatype indexed_type;
   MPI Init(&argc, &argv);
   MPI_Comm_rank(MPI_COMM_WORLD, &rank);
   /* Initialize data */
   for (int i = 0; i < size; i++) {</pre>
      data[i] = rank == 0 ? i : -1:
   /* Create indexed datatype */
   MPI_Type_indexed(3, block_lengths, displacements, MPI_DOUBLE, &indexed_type);
   MPI_Type_commit(&indexed_type);
```

Homogenous datatypes of variable strides: MPI_Type_indexed()

```
if (rank == 0) {
   /* Send data using indexed type */
   MPI Send(data, 1, indexed_type, 1, 0, MPI_COMM_WORLD);
} else if (rank == 1) {
   /* Receive data using indexed type */
   MPI Recv(data. 1. indexed_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
   printf("Rank %d: Received data: ", rank);
   for (int i = 0; i < size; i++) {
      printf("%.1f ", data[i]);
   printf("\n");
/* Clean up */
MPI Type free(&indexed type):
MPI Finalize():
return 0:
```

Rank 1: Received data: 0.0 1.0 -1.0 3.0 -1.0 -1.0 -1.0 7.0 8.0 9.0

Homogenous datatypes of variable strides:

MPI_Type_create_hindexed()



nb=4, blocks legths=(2,1,2,1), displacements=(2,10,14,24)

Figure: MPI_Type_create_hindexed() constructor

Homogenous datatypes of variable strides: MPI_Type_indexed()

• Example : triangular matrix

In the following example, each of the two processes:

- 1. Initializes its matrix (positive growing numbers on process 0 and negative decreasing numbers on process 1).
- 2. Constructs its datatype: triangular matrix (superior for the process 0 and inferior for the process 1).
- Sends its triangular matrix to the other process and receives back a triangular matrix which it stores in the same place which was occupied by the sent matrix.
- 4. Frees its resources.



Exchange between the two processes

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
int main(int argc,char *argv[]) {
   int rank, i, j;
   int n=8, tag=100,sign=1;
   float a[n][n];
   MPI_Datatype type_triangle;
   MPI_Status statut;
   int block_lengths[n],displacements[n];
   MPI_Init(&argc,&argv);
   MPI_Comm_rank(MPI_COMM_WORLD,&rank);
   /* Initialisation of the matrix on each process */
   if (rank == 1) sign=-1;
   for(i=0;i<n;i++)
      for(i=0:i<n:i++)
          a[i][j]=sign*(1+i*n+j);
```

```
/* Create type triangle for each process */
if (rank == 0) {
   for (i = 0: i < n: i++) {
      block_lengths[i] = i; // Upper triangle excluding diagonal
      displacements[i] = n * i; // Row offsets
} else {
   for (i = 0; i < n; i++) {
      block lengths[i] = n - i - 1; // Lower triangle excluding diagonal
      displacements[i] = (n + 1) * i + 1; // Row offsets
MPI_Type_indexed(n,block_lengths,displacements,MPI_FLOAT,&type_triangle);
MPI Type commit(&type triangle):
/* Swap of matrix */
MPI_Sendrecv_replace(a,1,type_triangle,(rank+1)%2,tag,(rank+1)%2,tag,MPI_COMM_WORLD,&statut);
/* Free the triangle datatype */
MPI Type free(&type triangle):
MPI Finalize():
```

```
mpirun -n 2 matrix exchange
Swapped matrix on rank 0:
1.0
      2.0
             3.0
                    4.0
                           5.0
                                   6.0
                                         7.0
                                                8.0
-2.0
                           13.0
                                  14.0
                                         15.0
      10.0
             11.0
                   12.0
                                                16.0
-3.0
      -4.0
            19.0
                    20.0
                           21.0
                                  22.0
                                         23.0
                                                24.0
-5.0
     -6.0
            -7.0
                    28.0
                           29.0
                                  30.0
                                         31.0
                                                32.0
            -12.0
                           37.0
-8.0
      -11.0
                   -13.0
                                  38.0
                                         39.0
                                                40.0
-14.0
     -15.0 -16.0 -20.0
                           -21.0
                                   46.0
                                          47.0
                                                 48.0
-22.0
      -23.0
             -24.0
                    -29.0
                           -30.0
                                   -31.0
                                          55.0
                                                 56.0
      -38.0
             -39.0
                    -40.0
                           -47.0
                                   -48.0
                                          -56.0
                                                 64 0
Swapped matrix on rank 1:
                    18.0
                           25.0
                                  26.0
                                         27.0
                                                33.0
            17.0
-9.0 -10.0
             34 0
                    35.0
                           36.0
                                  41.0
                                         42.0
                                                43.0
-17.0 -18.0
             -19.0 44.0
                           45.0
                                   49.0
                                          50.0
                                                 51.0
-25.0 -26.0 -27.0 -28.0
                           52.0
                                   53.0
                                          54.0
                                                 57.0
-33.0 -34.0 -35.0
                    -36.0
                           -37.0
                                   58.0
                                          59.0
                                                 60.0
-41.0 -42.0 -43.0
                    -44.0
                           -45.0
                                   -46.0
                                          61.0
                                                 62.0
-49.0 -50.0 -51.0 -52.0
                           -53.0
                                                 63.0
                                   -54.0
                                          -55.0
-57.0 -58.0 -59.0 -60.0 -61.0
                                  -62.0
                                         -63.0
                                                -64.0
```

Size of datatype: MPI_Type_size()

- In order to exchange derived datatypes in communications, it is important to obtain the size of these datatypes.
- The MPI_Type_size() subroutine provides the size of the derived datatype in bytes.

```
int MPI_Type_size(MPI_Datatype datatype,int *typesize)
```

The extent of a datatype is the memory space occupied by this datatype (in bytes).
 This value is used to calculate the position of the next datatype element (i.e. the stride between two successive datatype elements).

```
int MPI_Type_get_extent(MPI_Datatype datatype,MPI_Aint *lb, MPI_Aint *extent)
```

Heterogenous datatype

- MPI_Type_create_struct() call allows creating a set of data blocks indicating the type, the count and the displacement of each block.
- It is the most general datatype constructor. It further generalizes MPI_Type_indexed()
 by allowing a different datatype for each block.

Compute displacements

- MPI_Type_create_struct() is useful for creating MPI datatypes corresponding to Fortran derived datatypes or to C structures.
- The memory alignment of heterogeneous data structures is different for each architecture and each compiler.
- Warning, you have to check the extent of the MPI datatypes obtained.
- MPI_Get_address() provides the address of a variable. It's equivalent of & operator in C.
- Warning, even in C, it is better to use this subroutine for portability reasons.
- It's advised to use MPI_Aint_add() and MPI_Aint_diff() to add or substract addresses

```
int MPI_Get_address(const void *variable,MPI_Aint *address_variable)
MPI_Aint MPI_Aint_add(MPI_Aint base, MPI_Aint disp)
MPI_Aint MPI_Aint_diff(MPI_Aint addr1, MPI_Aint addr2)
```

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib h>
#include <stdbool b>
struct Particle {
   char category[5]:
   int mass:
   float coords[3]:
   bool class:
}:
int main(int argc,char *argv[]) {
   int rank.i:
   int n=1000,tag=100;
   int blocks_length[4];
   MPI_Datatype types[4],type_particle,temp;
   MPI Status statut:
   MPI Aint addresses[5].displacements[5].lb.extent:
   struct Particle p[n].temp p[n]:
   MPI Init(&argc,&argv):
   MPI Comm rank(MPI COMM WORLD, &rank):
```

```
/* Construction of the datatype */
types[0] = MPI_CHARACTER; types[1] = MPI_INT;
types[2] = MPI FLOAT: types[3] = MPI LOGICAL:
blocks length[0]=5:blocks length[1]=1:
blocks length[2]=3:blocks length[3]=1:
MPI Get address(&(p[0]),&(addresses[0])):
MPI_Get_address(&(p[0].category),&(addresses[1]));
MPI Get address(&(p[0].mass),&(addresses[2]));
MPI Get address(&(p[0].coords),&(addresses[3])):
MPI Get address(&(p[0].class),&(addresses[4]));
/* Compute displacements relative to start address */
for (i=0:i<4:i++) displacements[i] = MPI Aint diff(addresses[i+1].addresses[0]):</pre>
MPI Type create struct(4.blocks length.displacements.types.&temp):
MPI Get address(&(p[1]).&(addresses[1])):
1b=0:
extent = MPI Aint diff(addresses[1].addresses[0]);
MPI Type create resized(temp.lb.extent.&type particle):
/* Validation of type */
MPI Type commit(&type particle):
```

```
/* Initialization of particles on each process */
/* Send particles */
if (rank == 0) {
    MPI_Send(&(p[0]),n,type_particle,1,tag,MPI_COMM_WORLD);
} else {
    MPI_Recv(&(temp_p[0]),n,type_particle,0,tag,
    MPI_COMM_WORLD,&statut);
}

/* Free type */
MPI_Type_free(&type_particle);
MPI_Finalize();
}
```

Memento

Subroutines	Blocks_lengths	Strides	Old_types
MPI_Type_Contiguous() MPI_Type_[Create_H]Vector() MPI_Type_[Create_H]Indexed() MPI_Type_Create_Struct()	constant	constant	constant
	constant*	constant*	constant
	variable	variable	constant
	variable	variable	variable

(*) hidden parameter, equal to 1